



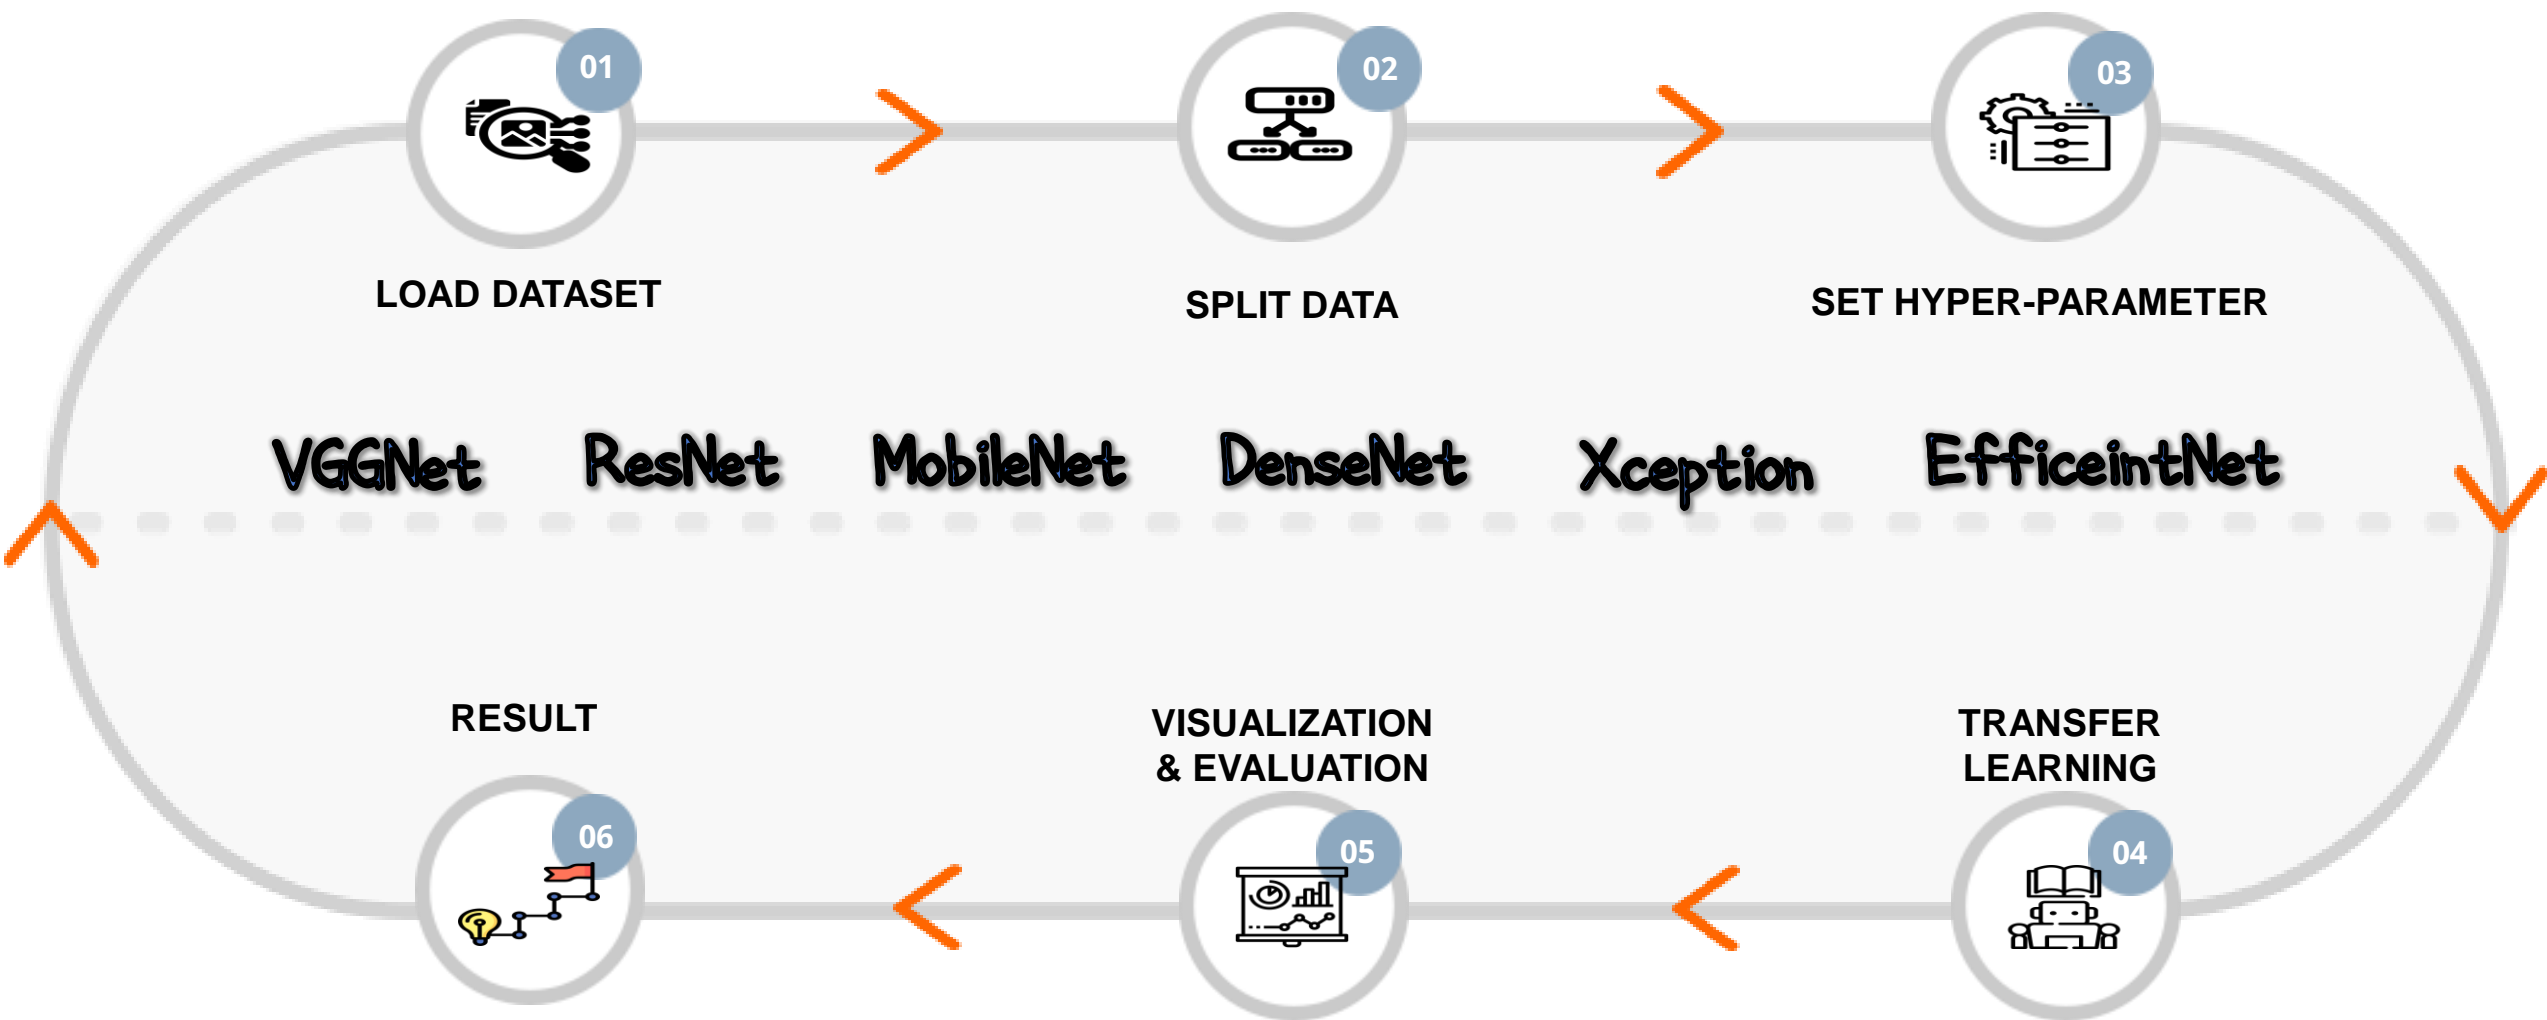
Leukemia Image Classification

- TRANSFER LEARNING
- HYPER PARAMETER TUNING



Outline

Outline





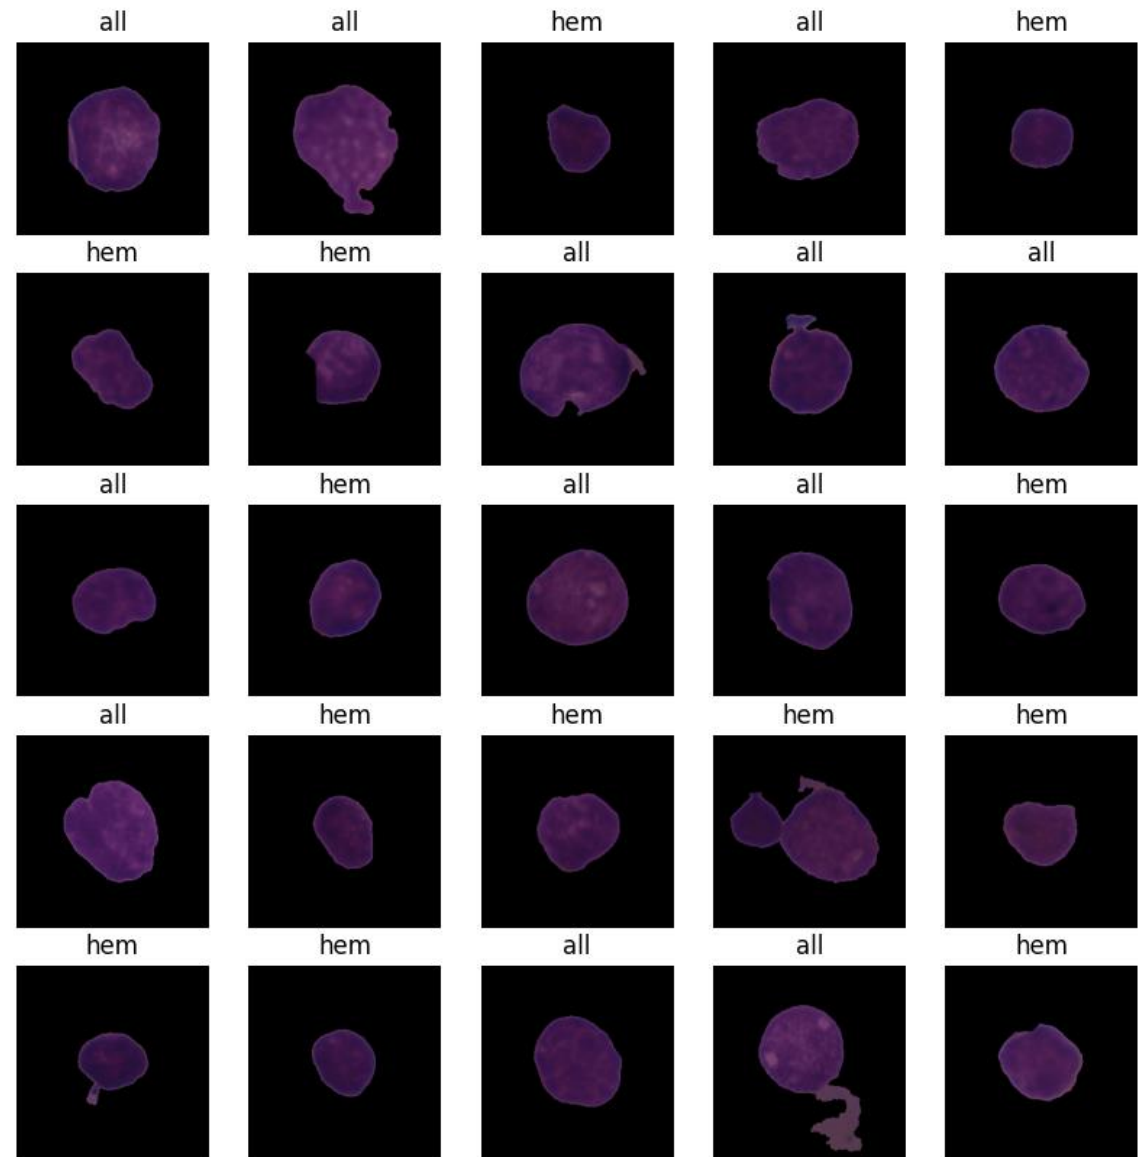
Dataset

1. DATASET :

LEUKEMIA IMAGE CLASSIFICATION

ALL VS HEM

- ALL : Leukemia blasts
- HEM : Normal Blood Cells



ALL (Leukemia blasts) VS HEM (Normal Blood Cells)



Preprocessing

Split Data from Training Dataset:

- Training data : 0.9
- Validation data : 0.05
- Test data : 0.05

```
[ ] # Split data
edir="/content/drive/MyDrive/Computer Vision/Image Classification/4.Leukemia Image Classification/C-NMC_Leukemia/training_data"
trsplit=.9
vsplit=.05
train_df, test_df, valid_df= preprocess(edir, trsplit, vsplit)

train_df length: 9594 test_df length: 534 valid_df length: 533
[6544, 3050]
```

Make train data set balance, resize the image

```
[ ] # The train data set is not balanced. To balance it use the trim function defined above to limit the maximum samples in a class to 3050 samples
max_samples=3050
min_samples=0
column='labels'
working_dir = r'./'
img_size=(224,224)
train_df=trim(train_df, max_samples, min_samples, column)
# The train_df dataframe is now balanced with 140 samples per class

Original Number of classes in dataframe: 2
[3050, 3050]
```

Create train, test and validation generators

```
# Set Hyperparameter
channels = 3
batch_size = 32
img_shape=(img_size[0], img_size[1], channels)
length=len(test_df)
test_batch_size=sorted([int(length/n) for n in range(1,length+1) if length % n ==0 and length/n<=80],reverse=True)[0]
test_steps=int(length/test_batch_size)
print ( 'test batch size: ',test_batch_size, ' test steps: ', test_steps)

# Set Train, Test, Valid generators
def scalar(img):
    return img # Network expects pixels in range 0 to 255 so no scaling is required
trgen=ImageDataGenerator(preprocessing_function=scalar, horizontal_flip=True)
tvgen=ImageDataGenerator(preprocessing_function=scalar)

msg='
for the train generator'
print(msg, '\n', end='')
train_gen=trgen.flow_from_dataframe( train_df,
                                     x_col='filepaths',
                                     y_col='labels',
                                     target_size=img_size,
                                     class_mode='categorical',
                                     color_mode='rgb',
                                     shuffle=True,
                                     batch_size=batch_size)

msg='
for the test generator'
print(msg, '\n', end='')
test_gen=tvgen.flow_from_dataframe( test_df,
                                   x_col='filepaths',
                                   y_col='labels',
                                   target_size=img_size,
                                   class_mode='categorical',
                                   color_mode='rgb',
                                   shuffle=False,
                                   batch_size=test_batch_size)

msg='
for the validation generator'
print(msg, '\n', end='')
valid_gen=tvgen.flow_from_dataframe( valid_df,
                                     x_col='filepaths',
                                     y_col='labels',
                                     target_size=img_size,
                                     class_mode='categorical',
                                     color_mode='rgb',
                                     shuffle=True,
                                     batch_size=batch_size)
```



```
test_batch_size: 6 test steps: 89
Found 6100 validated image filenames belonging to 2 classes.
Found 534 validated image filenames belonging to 2 classes.
Found 533 validated image filenames belonging to 2 classes.
```



Hyper-parameter Tuning

<u>Hyperparameter</u>	
Epoch	20
Optimizer	Adam
Loss	Binary Cross Entropy
Activation Function	Sigmoid, ReLU
Regularization	L1(0.006), L2(0.016)
Dropout	0.5
Batch Normalization	Momentum(0.99), Epsilon(0.001)
Learning Rate	0.00001

Hyper-parameter

Define Callbacks

```
[ ] class LRA(keras.callbacks.Callback):
    def __init__(self, model, base_model, patience, stop_patience, threshold, factor, dwell, batches, initial_epoch, epochs, ask_epoch):
        super(LRA, self).__init__()
        self.model=model
        self.base_model=base_model
        self.patience=patience # specifies how many epochs without improvement before learning rate is adjusted
        self.stop_patience=stop_patience # specifies how many times to adjust lr without improvement to stop training
        self.threshold=threshold # specifies training accuracy threshold when lr will be adjusted based on validation loss
        self.factor=factor # factor by which to reduce the learning rate
        self.dwell=dwell
        self.batches=batches # number of training batch to runn per epoch
        self.initial_epoch=initial_epoch
        self.epochs=epochs
        self.ask_epoch=ask_epoch
        self.ask_epoch_initial=ask_epoch # save this value to restore if restarting training

        # callback variables
        self.count=0 # how many times lr has been reduced without improvement
        self.stop_count=0
        self.best_epoch=1 # epoch with the lowest loss
        self.initial_lr=float(tf.keras.backend.get_value(model.optimizer.lr)) # get the initial learning rate and save it
        self.highest_tracc=0.0 # set highest training accuracy to 0 initially
        self.lowest_vloss=np.inf # set lowest validation loss to infinity initially
        self.best_weights=self.model.get_weights() # set best weights to model's initial weights
        self.initial_weights=self.model.get_weights() # save initial weights if they have to get restored

# Set Callback Option
patience= 10
stop_patience =3
threshold=.9
factor=.5
dwell=True
freeze=False
ask_epoch=20
batches=train_steps

callbacks=[LRA(model=model6,base_model= base_model, patience=patience, stop_patience=stop_patience, threshold=threshold,
               factor=factor,dwell=dwell, batches=batches, initial_epoch=0, epochs=epochs, ask_epoch=ask_epoch )]
```

Callbacks



Model Summary

2. Model Summary

Model: "VGGNet19"

Layer (type)	Output Shape	Param #
input_4 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv4 (Conv2D)	(None, 56, 56, 256)	590080
...		
Total params: 20,092,354		
Trainable params: 20,091,330		
Non-trainable params: 1,024		

VGGNet : VGGNet19

Model: "ResNet50"

Layer (type)	Output Shape	Param #	Connected to
input_3 (InputLayer)	[(None, 224, 224, 3)]	0	[]
conv1_pad (ZeroPadding2D)	(None, 230, 230, 3)	0	['input_3[0][0]']
conv1_conv (Conv2D)	(None, 112, 112, 64)	9472	['conv1_pad[0][0]']
conv1_bn (BatchNormalization)	(None, 112, 112, 64)	256	['conv1_conv[0][0]']
conv1_relu (Activation)	(None, 112, 112, 64)	0	['conv1_bn[0][0]']
pool1_pad (ZeroPadding2D)	(None, 114, 114, 64)	0	['conv1_relu[0][0]']
pool1_pool (MaxPooling2D)	(None, 56, 56, 64)	0	['pool1_pad[0][0]']
conv2_block1_1_conv (Conv2D)	(None, 56, 56, 64)	4160	['pool1_pool[0][0]']
...			
Total params: 23,858,434			
Trainable params: 23,801,218			
Non-trainable params: 57,216			

ResNet : ResNet50

2. Model Summary

Model: "MobileNet"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
conv1 (Conv2D)	(None, 112, 112, 32)	864
conv1_bn (BatchNormalization)	(None, 112, 112, 32)	128
conv1_relu (ReLU)	(None, 112, 112, 32)	0
conv_dw_1 (DepthwiseConv2D)	(None, 112, 112, 32)	288
conv_dw_1_bn (BatchNormalization)	(None, 112, 112, 32)	128
conv_dw_1_relu (ReLU)	(None, 112, 112, 32)	0
conv_pw_1 (Conv2D)	(None, 112, 112, 64)	2048
conv_pw_1_bn (BatchNormalization)	(None, 112, 112, 64)	256
...		
Total params: 3,364,418		
Trainable params: 3,340,482		
Non-trainable params: 23,936		

MobileNet : MobileNet

Model: "DenseNet121"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 224, 224, 3)]	0	['input_1[0][0]']
zero_padding2d (ZeroPadding2D)	(None, 230, 230, 3)	0	['input_1[0][0]']
conv1/conv (Conv2D)	(None, 112, 112, 64)	9408	['zero_padding2d[0][0]']
conv1/bn (BatchNormalization)	(None, 112, 112, 64)	256	['conv1/conv[0][0]']
conv1/relu (Activation)	(None, 112, 112, 64)	0	['conv1/bn[0][0]']
zero_padding2d_1 (ZeroPadding2D)	(None, 114, 114, 64)	0	['conv1/relu[0][0]']
pool1 (MaxPooling2D)	(None, 56, 56, 64)	0	['zero_padding2d_1[0][0]']
conv2_block1_0_bn (BatchNormalization)	(None, 56, 56, 64)	256	['pool1[0][0]']
...			
Total params: 7,173,058			
Trainable params: 7,087,362			
Non-trainable params: 85,696			

DenseNet : DenseNet121

2. Model Summary

Model: "Xception"

Layer (type)	Output Shape	Param #	Connected to
=====			
input_5 (InputLayer)	(None, 224, 224, 3)	0	[]
block1_conv1 (Conv2D)	(None, 111, 111, 32)	864	['input_5[0][0]']
block1_conv1_bn (BatchNormaliz ation)	(None, 111, 111, 32)	128	['block1_conv1[0][0]']
block1_conv1_act (Activation)	(None, 111, 111, 32)	0	['block1_conv1_bn[0][0]']
block1_conv2 (Conv2D)	(None, 109, 109, 64)	18432	['block1_conv1_act[0][0]']
block1_conv2_bn (BatchNormaliz ation)	(None, 109, 109, 64)	256	['block1_conv2[0][0]']
block1_conv2_act (Activation)	(None, 109, 109, 64)	0	['block1_conv2_bn[0][0]']
...			
Total params: 21,132,202			
Trainable params: 21,073,578			
Non-trainable params: 58,624			

Xception : Xception

Model: "EfficientNetB1"


Layer (type)	Output Shape	Param #	Connected to
=====			
input_2 (InputLayer)	(None, 224, 224, 3)	0	[]
rescaling (Rescaling)	(None, 224, 224, 3)	0	['input_2[0][0]']
normalization (Normalization)	(None, 224, 224, 3)	7	['rescaling[0][0]']
rescaling_1 (Rescaling)	(None, 224, 224, 3)	0	['normalization[0][0]']
stem_conv_pad (ZeroPadding2D)	(None, 225, 225, 3)	0	['rescaling_1[0][0]']
stem_conv (Conv2D)	(None, 112, 112, 32)	864	['stem_conv_pad[0][0]']
stem_bn (BatchNormalization)	(None, 112, 112, 32)	128	['stem_conv[0][0]']
stem_activation (Activation)	(None, 112, 112, 32)	0	['stem_bn[0][0]']
...			
block1a_dwconv (DepthwiseConv2	(None, 112, 112, 32)	288	['stem_activation[0][0]']
...			
Total params: 6,744,585			
Trainable params: 6,679,970			
Non-trainable params: 64,615			

EfficientNet : EfficientNetB0

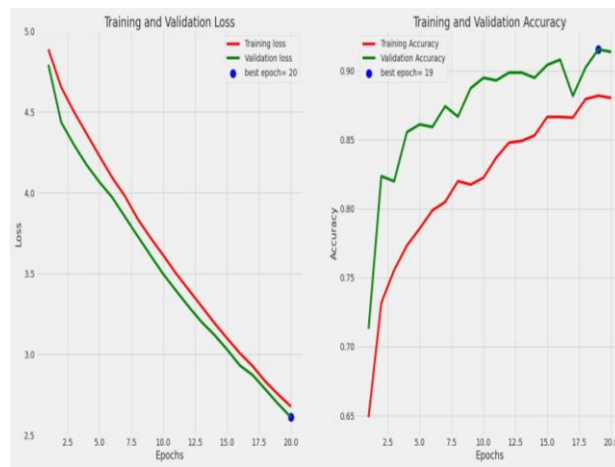


Results


3. Results : Loss and Accuracy



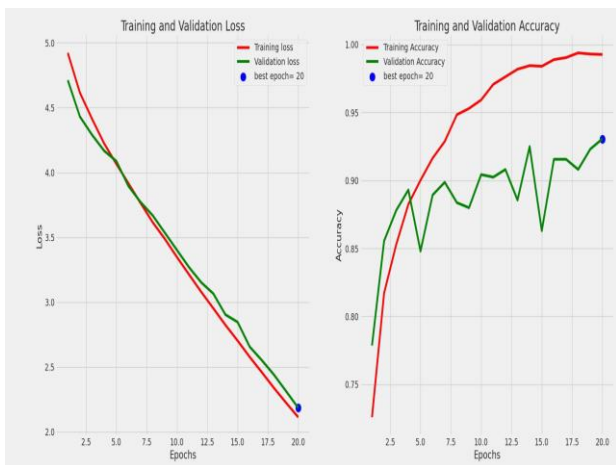
DenseNet121



EfficientNetB1



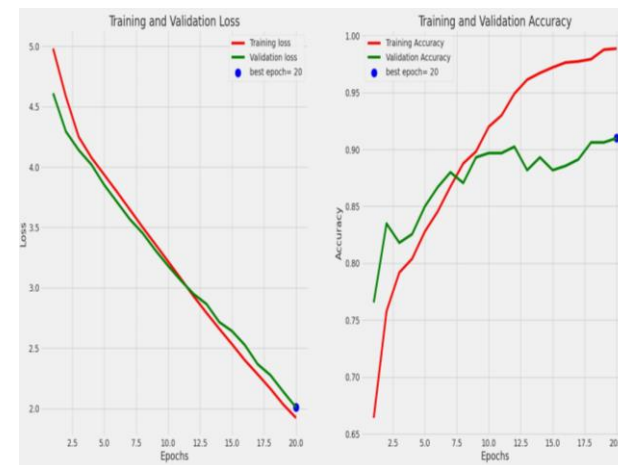
MobileNet



ResNet50



VGG19



Xception

3. Results : Classification Report

Classification Report:

	precision	recall	f1-score	support
all	0.9394	0.9368	0.9381	364
hem	0.8655	0.8706	0.8680	170
accuracy			0.9157	534
macro avg	0.9024	0.9037	0.9031	534
weighted avg	0.9159	0.9157	0.9158	534

DenseNet121

Classification Report:

	precision	recall	f1-score	support
all	0.9345	0.9011	0.9175	364
hem	0.8033	0.8647	0.8329	170
accuracy			0.8895	534
macro avg	0.8689	0.8829	0.8752	534
weighted avg	0.8927	0.8895	0.8905	534

EfficientNetB1

Classification Report:

	precision	recall	f1-score	support
all	0.9032	0.9231	0.9130	364
hem	0.8272	0.7882	0.8072	170
accuracy			0.8801	534
macro avg	0.8652	0.8557	0.8601	534
weighted avg	0.8790	0.8801	0.8794	534

MobileNet

Classification Report:

	precision	recall	f1-score	support
all	0.9339	0.9698	0.9515	364
hem	0.9295	0.8529	0.8896	170
accuracy			0.9326	534
macro avg	0.9317	0.9114	0.9205	534
weighted avg	0.9325	0.9326	0.9318	534

ResNet50

Classification Report:

	precision	recall	f1-score	support
all	0.9489	0.9176	0.9330	364
hem	0.8352	0.8941	0.8636	170
accuracy			0.9101	534
macro avg	0.8920	0.9059	0.8983	534
weighted avg	0.9127	0.9101	0.9109	534

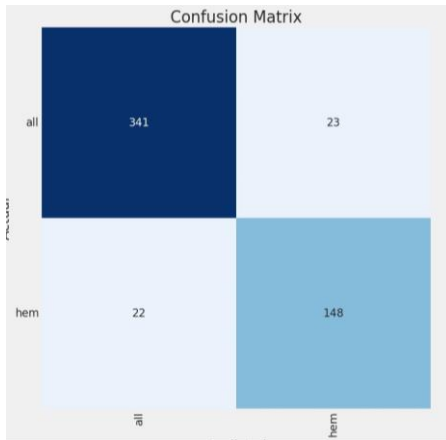
VGG19

Classification Report:

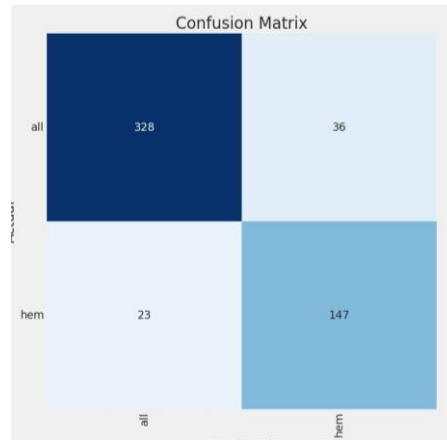
	precision	recall	f1-score	support
all	0.9096	0.9396	0.9243	364
hem	0.8608	0.8000	0.8293	170
accuracy			0.8951	534
macro avg	0.8852	0.8698	0.8768	534
weighted avg	0.8940	0.8951	0.8941	534

Xception

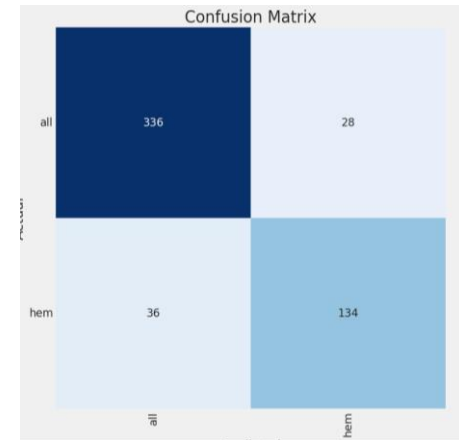
3. Results : Confusion Matrix




DenseNet121



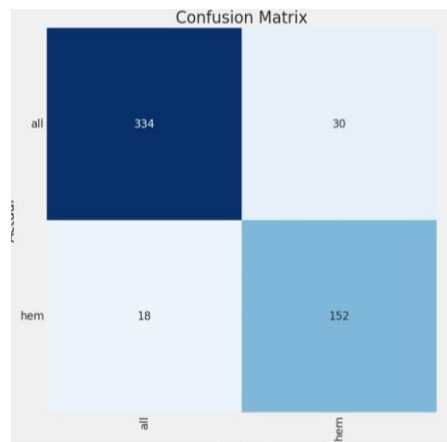
EfficientNetB1




MobileNet



ResNet50



VGG19



Xception

3. Results : Number of Errors

There were 45 errors in 534 test cases Model accuracy= 91.57 %



DenseNet121

There were 59 errors in 534 test cases Model accuracy= 88.95 %



EfficientNetB1

There were 64 errors in 534 test cases Model accuracy= 88.01 %



MobileNet

There were 36 errors in 534 test cases Model accuracy= 93.26 %




ResNet50

There were 48 errors in 534 test cases Model accuracy= 91.01 %



VGG19

There were 56 errors in 534 test cases Model accuracy= 89.51 %



Xception



Analysis

4. Analyzation : Loss and Accuracy

- Result Analysis 1 : Loss and Accuracy
- Training Loss and Accuracy, Evaluation Loss and Accuracy for each model
- Top 3 Models : [ResNet50](#), [DenseNet121](#), [VGGNet19](#)

	Train_Loss	Eval_Loss	Train_Accuracy	Eval_Accuracy
DenseNet121	1.6205	1.7057	0.9867	0.9287
MobileNet	2.6480	2.6754	0.8967	0.8668
EfficientNetB1	2.6775	2.6113	0.8803	0.9137
ResNet50	2.1126	2.1845	0.9926	0.9306
VGGNet19	1.6934	1.6692	0.9377	0.9250
Xception	1.9231	2.0096	0.9887	0.9099

Results : Loss and Accuracy


- Result Analysis 2 : Precision, Recall, F1-Score
- Precision, Recall, F1-Score of ALL vs Precision, Recall, F1-Score of HEM
- Top 3 Models : [ResNet50](#), [DenseNet121](#), [VGGNet19](#)

	Precision(ALL)	Recall(ALL)	F1-Score(ALL)	Precision(HEM)	Recall(HEM)	F1-Score(HEM)
DenseNet121	0.9394	0.9368	0.9381	0.8655	0.8706	0.8680
MobileNet	0.9032	0.9231	0.9130	0.8272	0.7882	0.8072
EfficientNetB1	0.9345	0.9011	0.9175	0.8033	0.8647	0.8329
ResNet50	0.9339	0.9698	0.9515	0.9295	0.8529	0.8896
VGGNet19	0.9489	0.9176	0.9330	0.8352	0.8941	0.8636
Xception	0.9096	0.9396	0.9243	0.8608	0.8000	0.8293

Results : Precision, Recall, F1-score (ALL & HEM)


4. Analyzation : Loss and Accuracy

- Result Analysis 1 : Loss and Accuracy
- Top 3 Models : ResNet50, DenseNet121, VGGNet19



Results : Loss and Accuracy

- Result Analysis 2 : Precision, Recall, F1-Score
- Top 3 Models : ResNet50, DenseNet121, VGGNet19



Results : Precision, Recall, F1-score



THANK YOU

Contacts

- E-mail : johnnyworld9278@gmail.com
- GitHub : <https://github.com/jpjp92>